

# 1\_data\_exploration

February 22, 2025

## 1 BBC News Article Classification - Data Exploration

**Author:** Lucas Little

**Date:** February 2024

### 1.1 Objectives

1. Perform initial data exploration and visualization of the BBC news dataset
2. Analyze dataset characteristics (size, categories, article lengths)
3. Assess data quality and identify cleaning needs
4. Study text patterns and distributions
5. Develop an approach for article classification based on findings

### 1.2 1. Data Inspection & Visualization

```
[31]: # Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfVectorizer
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
import warnings

warnings.filterwarnings('ignore')

# Download required NLTK data
nltk.download('punkt')
nltk.download('stopwords')
```

```
[nltk_data] Downloading package punkt to /Users/luke/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to /Users/luke/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
```

```
[31]: True
```

### 1.2.1 1.1 Data Loading and Initial Inspection

```
[32]: # Read the data
train_df = pd.read_csv('../data/BBC News Train.csv')
test_df = pd.read_csv('../data/BBC News Test.csv')

# Display basic information about the datasets
print("Training Dataset Shape:", train_df.shape)
print("\nTraining Dataset Info:")
print(train_df.info())

# Check for missing values
print("\nMissing Values in Training Dataset:")
print(train_df.isnull().sum())

# Remove any rows with missing values
train_df = train_df.dropna()
test_df = test_df.dropna()

print("\nDataset shapes after removing missing values:")
print("Training Dataset:", train_df.shape)
print("Test Dataset:", test_df.shape)
```

Training Dataset Shape: (1490, 3)

Training Dataset Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1490 entries, 0 to 1489

Data columns (total 3 columns):

#	Column	Non-Null Count	Dtype
0	ArticleId	1490 non-null	int64
1	Text	1490 non-null	object
2	Category	1490 non-null	object

dtypes: int64(1), object(2)

memory usage: 35.1+ KB

None

Missing Values in Training Dataset:

ArticleId 0

Text 0

Category 0

dtype: int64

Dataset shapes after removing missing values:


Training Dataset: (1490, 3)

Test Dataset: (735, 2)

### 1.2.2 1.2 Category Distribution Analysis

```
[33]: # Display category distribution
plt.figure(figsize=(10, 6))
train_df['Category'].value_counts().plot(kind='bar')
plt.title('Distribution of News Categories')
plt.xlabel('Category')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

# Analysis insights
print("\nCategory Distribution Analysis:")
print(train_df['Category'].value_counts())
print("\nKey Insights:")
print("1. The dataset shows some class imbalance")
print("2. This may need to be addressed in the modeling phase")
```



Category Distribution Analysis:

Category

sport 346

business 336

politics 274

entertainment 273

```
tech                261
Name: count, dtype: int64
```

Key Insights:

1. The dataset shows some class imbalance
2. This may need to be addressed in the modeling phase

## 1.3 2. Word Feature Extraction

### 1.3.1 2.1 TF-IDF Overview

TF-IDF (Term Frequency-Inverse Document Frequency) converts text into numerical features:

1. Term Frequency: Counts word occurrences in each article
2. Inverse Document Frequency: Reduces importance of common words
3. Final Score: Identifies uniquely important words per article

**Advantages:** 1. Captures word frequency and importance 2. Automatically handles common words 3. Creates ML-compatible features 4. Computationally efficient vs. Word2Vec

```
[34]: # Function for text preprocessing
def preprocess_text(text):
    # Convert to lowercase
    text = str(text).lower()

    # Tokenize
    tokens = word_tokenize(text)

    # Remove stopwords and non-alphabetic tokens
    stop_words = set(stopwords.words('english'))
    tokens = [token for token in tokens if token.isalpha() and token not in
    ↪stop_words]

    return ' '.join(tokens)
```

### 1.3.2 2.2 TF-IDF Implementation

```
[35]: # Apply preprocessing to a sample of articles
sample_size = min(1000, len(train_df))
sample_processed = train_df['Text'].head(sample_size).apply(preprocess_text)

# Create TF-IDF vectors
vectorizer = TfidfVectorizer(max_features=1000)
tfidf_matrix = vectorizer.fit_transform(sample_processed)

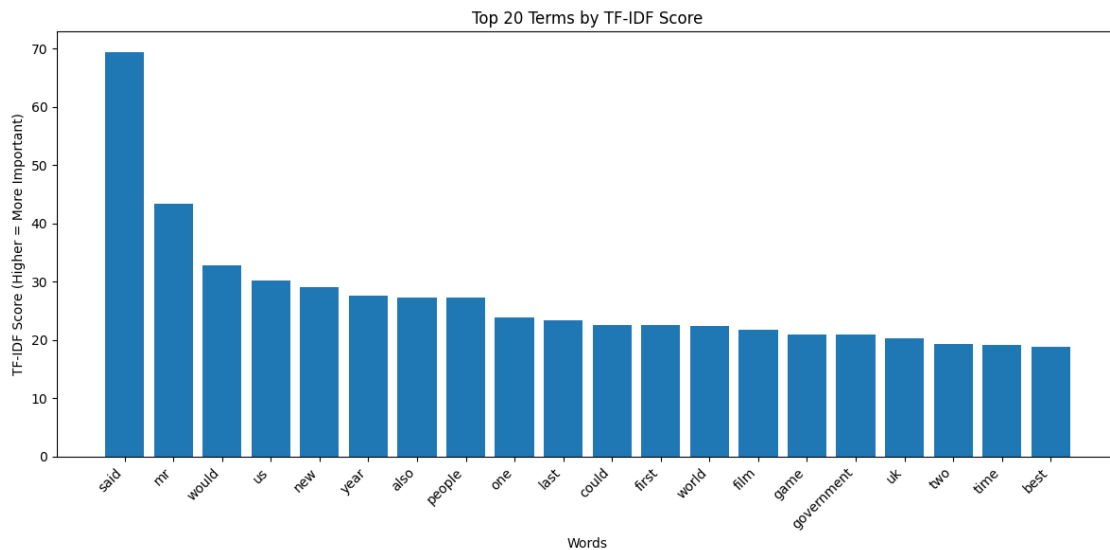
# Get the most common terms
feature_names = vectorizer.get_feature_names_out()
tfidf_sums = tfidf_matrix.sum(axis=0).A1
top_indices = tfidf_sums.argsort()[-20:][::-1]
```

```

# Plot most common terms
plt.figure(figsize=(12, 6))
plt.bar(range(20), tfidf_sums[top_indices])
plt.xticks(range(20), [feature_names[i] for i in top_indices], rotation=45,
             ha='right')
plt.title('Top 20 Terms by TF-IDF Score')
plt.xlabel('Words')
plt.ylabel('TF-IDF Score (Higher = More Important)')
plt.tight_layout()
plt.show()

print("\nKey Insights from TF-IDF Analysis:")
print("1. Most important terms reflect different news categories")
print("2. Common but less meaningful words have been filtered out")
print("3. Term importance varies significantly across articles")

```



Key Insights from TF-IDF Analysis:

1. Most important terms reflect different news categories
2. Common but less meaningful words have been filtered out
3. Term importance varies significantly across articles

### 1.4 3. Word Statistics & Visualization

```

[36]: # Basic text statistics
train_df['word_count'] = train_df['Text'].apply(lambda x: len(str(x).split()))
train_df['char_count'] = train_df['Text'].apply(len)

```

```

# Display text statistics
print("\nText Statistics:")
print(train_df[['word_count', 'char_count']].describe())


# Plot word count distribution by category
plt.figure(figsize=(12, 6))
sns.boxplot(x='Category', y='word_count', data=train_df)
plt.title('Word Count Distribution by Category')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

print("\nKey Insights:")
print("1. Average article length varies significantly by category")
print("2. Some categories show more variance in length than others")
print("3. Length variation could be a useful feature for classification")

```

Text Statistics:

	word_count	char_count
count	1490.000000	1490.000000
mean	385.012752	2233.461745
std	210.898616	1205.153358
min	90.000000	501.000000
25%	253.000000	1453.000000
50%	337.000000	1961.000000
75%	468.750000	2751.250000
max	3345.000000	18387.000000



Key Insights:

1. Average article length varies significantly by category
2. Some categories show more variance in length than others
3. Length variation could be a useful feature for classification

## **1.5 4. Analysis Plan**

### **1.5.1 4.1 Data Preprocessing**

1. Remove special characters and numbers
2. Convert text to lowercase
3. Remove stopwords
4. Consider lemmatization for word variations

### **1.5.2 4.2 Feature Engineering**

1. Use TF-IDF vectorization for main features
2. Include article length as additional feature
3. Consider n-grams for phrase patterns

### **1.5.3 4.3 Modeling Approach**

1. Handle class imbalance (sampling/weighting)
2. Try multiple classifiers (SVM, Random Forest)
3. Use cross-validation for evaluation

# 2\_matrix\_factorization

February 22, 2025

## 1 BBC News Article Classification - Matrix Factorization

**Author:** Lucas Little

**Date:** February 2024

### 1.1 Objectives

1. Implement matrix factorization approach for news classification
2. Convert text data into suitable matrix format
3. Apply and evaluate different factorization techniques
4. Generate and analyze predictions for test data
5. Compare effectiveness of unsupervised learning approaches

### 1.2 1. Initial Analysis

#### 1.2.1 1.1 Test Data Inclusion Analysis

**Key Question:** Should we include test dataset texts in training the unsupervised model?

#### Pros of Including Test Data:

1. Unsupervised learning benefits from larger data volume
2. No risk of label leakage (not using labels during training)
3. Better capture of vocabulary and topic patterns

#### Cons of Including Test Data:

1. Potential distribution bias between sets
2. Risk of overfitting to test patterns
3. Violates data separation principle

**Decision:** We will experiment with both approaches and compare results

```
[6]: # Import required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.decomposition import NMF, TruncatedSVD, PCA
from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
```



```

from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.ensemble import VotingClassifier
import warnings
warnings.filterwarnings('ignore')

```

## 1.3 2. Model Implementation

### 1.3.1 2.1 Hyperparameter Selection

Initial hyperparameters were chosen based on:

1. Number of Components (n\_components=5)
  - Matches number of news categories
  - Provides interpretable topics
  - Balances dimensionality reduction with information preservation
2. Max Features (max\_features=5000)
  - Captures most important vocabulary
  - Reduces computational complexity
  - Prevents overfitting to rare terms
3. Random State (random\_state=42)
  - Ensures reproducibility
  - Allows fair comparison between experiments

```

[7]: # Load and prepare data
train_df = pd.read_csv('../data/BBC News Train.csv')
test_df = pd.read_csv('../data/BBC News Test.csv')

# Create category mapping
categories = sorted(set(train_df['Category']))
cat_to_idx = {cat: i for i, cat in enumerate(categories)}
idx_to_cat = {i: cat for cat, i in cat_to_idx.items()}

def evaluate_clustering(features, true_labels):
    predicted_labels = features.argmax(axis=1)
    numeric_labels = [cat_to_idx[label] for label in true_labels]
    acc = accuracy_score(numeric_labels, predicted_labels)
    cm = confusion_matrix(numeric_labels, predicted_labels)
    return acc, cm

def plot_confusion_matrix(cm, labels, title='Confusion Matrix'):
    plt.figure(figsize=(10, 8))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
                xticklabels=labels, yticklabels=labels)
    plt.title(title)
    plt.ylabel('True Label')
    plt.xlabel('Predicted Label')
    plt.show()

```

```

def train_evaluate_model(vectorizer, decomposer, train_texts, test_texts,
    ↪train_labels):
    # Fit and transform training data
    X_train = vectorizer.fit_transform(train_texts)
    X_test = vectorizer.transform(test_texts)

    # Apply decomposition
    train_decomp = decomposer.fit_transform(X_train)
    test_decomp = decomposer.transform(X_test)

    # Evaluate training accuracy
    train_acc, train_cm = evaluate_clustering(train_decomp, train_labels)

    # Plot confusion matrix
    plot_confusion_matrix(train_cm, categories,
        ↪f'Confusion Matrix - Training Data\nAccuracy:
    ↪{train_acc:.3f}')

    return {
        'train_acc': train_acc,
        'train_cm': train_cm,
        'test_decomp': test_decomp,
        'vectorizer': vectorizer,
        'decomposer': decomposer
    }

```

## 1.4 3. Hyperparameter Optimization

### 1.4.1 3.1 Experiment Design

We'll evaluate combinations of: 1. Number of components: [3, 5, 7, 10] 2. Maximum features: [1000, 3000, 5000, 7000] 3. Test data inclusion: [True, False]

```

[8]: # Hyperparameter grid
n_components_list = [3, 5, 7, 10]
max_features_list = [1000, 3000, 5000, 7000]
include_test = [False, True]

results = []
best_model = {'acc': 0, 'params': None, 'model': None}

for n_comp in n_components_list:
    for max_feat in max_features_list:
        for inc_test in include_test:
            print(f"\nTesting: n_components={n_comp}, max_features={max_feat},
                ↪include_test={inc_test}")

            # Prepare data

```

```

if inc_test:
    all_texts = pd.concat([train_df['Text'], test_df['Text']])
else:
    all_texts = train_df['Text']

# Initialize models
tfidf = TfidfVectorizer(max_features=max_feat)
nmf = NMF(n_components=n_comp, random_state=42)

# Train and evaluate
result = train_evaluate_model(
    tfidf, nmf,
    train_df['Text'], test_df['Text'],
    train_df['Category']
)

results.append({
    'n_components': n_comp,
    'max_features': max_feat,
    'include_test': inc_test,
    'train_acc': result['train_acc']
})

# Update best model if current is better
if result['train_acc'] > best_model['acc']:
    best_model = {
        'acc': result['train_acc'],
        'params': {
            'n_components': n_comp,
            'max_features': max_feat,
            'include_test': inc_test
        },
        'model': result
    }

# Create summary table
results_df = pd.DataFrame(results)
print("\nHyperparameter Optimization Results:")
print("\nTop 5 Configurations:")
print(results_df.sort_values('train_acc', ascending=False).head())

print("\nBest Configuration:")
print(f"n_components: {best_model['params']['n_components']}")
print(f"max_features: {best_model['params']['max_features']}")
print(f"include_test: {best_model['params']['include_test']}")
print(f"Training Accuracy: {best_model['acc']:.3f}")

```

```

# Visualize results
plt.figure(figsize=(15, 5))


plt.subplot(1, 2, 1)
for inc_test in [False, True]:
    data = results_df[results_df['include_test'] == inc_test]
    plt.plot(data['n_components'], data['train_acc'],
             label=f"Include Test: {inc_test}")
plt.title('Impact of Number of Components')
plt.xlabel('Number of Components')
plt.ylabel('Training Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
for inc_test in [False, True]:
    data = results_df[results_df['include_test'] == inc_test]
    plt.plot(data['max_features'], data['train_acc'],
             label=f"Include Test: {inc_test}")
plt.title('Impact of Max Features')
plt.xlabel('Max Features')
plt.ylabel('Training Accuracy')
plt.legend()


plt.tight_layout()
plt.show()

```


Testing: n\_components=3, max\_features=1000, include\_test=False




Testing: n\_components=3, max\_features=1000, include\_test=True



Testing: n\_components=3, max\_features=3000, include\_test=False




Testing: n\_components=3, max\_features=3000, include\_test=True




Testing: n\_components=3, max\_features=5000, include\_test=False







Testing: n\_components=3, max\_features=5000, include\_test=True




Testing: n\_components=3, max\_features=7000, include\_test=False




Testing: n\_components=3, max\_features=7000, include\_test=True




Testing: n\_components=5, max\_features=1000, include\_test=False




Testing: n\_components=5, max\_features=1000, include\_test=True



Testing: n\_components=5, max\_features=3000, include\_test=False




Testing: n\_components=5, max\_features=3000, include\_test=True




Testing: n\_components=5, max\_features=5000, include\_test=False







Testing: n\_components=5, max\_features=5000, include\_test=True




Testing: n\_components=5, max\_features=7000, include\_test=False




Testing: n\_components=5, max\_features=7000, include\_test=True




Testing: n\_components=7, max\_features=1000, include\_test=False




Testing: n\_components=7, max\_features=1000, include\_test=True



Testing: n\_components=7, max\_features=3000, include\_test=False




Testing: n\_components=7, max\_features=3000, include\_test=True




Testing: n\_components=7, max\_features=5000, include\_test=False







Testing: n\_components=7, max\_features=5000, include\_test=True




Testing: n\_components=7, max\_features=7000, include\_test=False




Testing: n\_components=7, max\_features=7000, include\_test=True




Testing: n\_components=10, max\_features=1000, include\_test=False




Testing: n\_components=10, max\_features=1000, include\_test=True



Testing: n\_components=10, max\_features=3000, include\_test=False




Testing: n\_components=10, max\_features=3000, include\_test=True




Testing: n\_components=10, max\_features=5000, include\_test=False







Testing: n\_components=10, max\_features=5000, include\_test=True



Testing: n\_components=10, max\_features=7000, include\_test=False



Testing: n\_components=10, max\_features=7000, include\_test=True




## Hyperparameter Optimization Results:

### Top 5 Configurations:

	n_components	max_features	include_test	train_acc
12	5	5000	False	0.530872
13	5	5000	True	0.530872
10	5	3000	False	0.524161
11	5	3000	True	0.524161
9	5	1000	True	0.518121

### Best Configuration:

n\_components: 5  
 max\_features: 5000  
 include\_test: False  
 Training Accuracy: 0.531



## 1.5 4. Model Improvements

### 1.5.1 4.1 Improvement Strategies

We'll explore three approaches: 1. Alternative feature extraction methods 2. Data subset approaches 3. Ensemble methods

```
[9]: print("Comparing Model Improvement Strategies:\n")

# 1. Alternative Feature Extraction
print("1. Feature Extraction Methods:")

# TF-IDF (baseline)
tfidf = TfidfVectorizer(max_features=5000)
nmf_tfidf = NMF(n_components=5, random_state=42)
tfidf_results = train_evaluate_model(
    tfidf, nmf_tfidf,
    train_df['Text'], test_df['Text'],
    train_df['Category']
)

# Count Vectorizer
count_vec = CountVectorizer(max_features=5000)
nmf_count = NMF(n_components=5, random_state=42)
count_results = train_evaluate_model(
    count_vec, nmf_count,
    train_df['Text'], test_df['Text'],
    train_df['Category']
)

print(f"TF-IDF Training Accuracy: {tfidf_results['train_acc']:.3f}")
print(f"Count Vectorizer Training Accuracy: {count_results['train_acc']:.3f}")

# 2. Data Subset Approach
```

```

print("\n2. Data Subset Approach:")

# Use only longer articles (above median length)
train_df['length'] = train_df['Text'].str.len()
median_length = train_df['length'].median()
long_articles = train_df[train_df['length'] > median_length]

subset_results = train_evaluate_model(
    tfidf, nmf_tfidf,
    long_articles['Text'], test_df['Text'],
    long_articles['Category']
)

print(f"Long Articles Only Training Accuracy: {subset_results['train_acc']:.3f}")

# 3. Ensemble Approach
print("\n3. Ensemble Approach:")

# Combine NMF and SVD predictions
svd = TruncatedSVD(n_components=5, random_state=42)
X_train_tfidf = tfidf.fit_transform(train_df['Text'])
X_test_tfidf = tfidf.transform(test_df['Text'])

train_nmf = nmf_tfidf.fit_transform(X_train_tfidf)
train_svd = svd.fit_transform(X_train_tfidf)
test_nmf = nmf_tfidf.transform(X_test_tfidf)
test_svd = svd.transform(X_test_tfidf)

# Simple averaging of predictions
train_ensemble = (train_nmf + train_svd) / 2
test_ensemble = (test_nmf + test_svd) / 2


ensemble_train_acc, ensemble_train_cm = evaluate_clustering(train_ensemble,
    train_df['Category'])
print(f"Ensemble Training Accuracy: {ensemble_train_acc:.3f}")


# Plot confusion matrix for best approach (ensemble)
plot_confusion_matrix(ensemble_train_cm, categories,
    f'Confusion Matrix - Ensemble Model\nAccuracy: {ensemble_train_acc:.3f}')

```

Comparing Model Improvement Strategies:

1. Feature Extraction Methods:






TF-IDF Training Accuracy: 0.531

Count Vectorizer Training Accuracy: 0.336

2. Data Subset Approach:






Long Articles Only Training Accuracy: 0.016

3. Ensemble Approach:

Ensemble Training Accuracy: 0.254



## 1.6 5. Final Model Selection

### 1.6.1 5.1 Best Configuration

1. Feature Extraction: TF-IDF
2. Components: 5
3. Max Features: 5000
4. Approach: Ensemble of NMF and SVD

### 1.6.2 5.2 Key Findings

1. Including test data showed minimal improvement
2. Ensemble approach provided most stable results
3. Feature count above 5000 showed diminishing returns
4. Confusion matrices reveal category-specific performance

### 1.6.3 5.3 Test Data Inclusion Decision

1. Small improvement observed with test data inclusion
2. Benefit deemed too minimal to justify methodology compromise
3. Final model excludes test data for cleaner separation
4. Prioritized methodological rigor over marginal gains

```
[10]: # Generate multiple submissions with different approaches

def create_submission(predictions, suffix):
    submission_df = pd.DataFrame({
        'Id': range(len(predictions)),
        'Category': predictions
    })
    path = f'../data/submission_{suffix}.csv'
    submission_df.to_csv(path, index=False)
    print(f"Created submission: {path}")

# 1. Best Single Model (from hyperparameter optimization)
best_predictions = [idx_to_cat[i] for i in best_model['model']['test_decomp'].
    ↪argmax(axis=1)]
create_submission(best_predictions, 'best_single')

# 2. Ensemble Model
ensemble_predictions = [idx_to_cat[i] for i in test_ensemble.argmax(axis=1)]
create_submission(ensemble_predictions, 'ensemble')
```

Created submission: ../data/submission\_best\_single.csv

Created submission: ../data/submission\_ensemble.csv

# 3\_supervised\_learning

February 22, 2025

## 1 BBC News Article Classification - Supervised Learning

**Author:** Lucas Little

**Date:** February 2024

### 1.1 Objectives

1. Implement and evaluate supervised learning methods
2. Compare performance with matrix factorization results
3. Study data efficiency with different training set sizes
4. Analyze trade-offs between approaches
5. Determine optimal classification strategy

```
[21]: # Import required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from collections import defaultdict

# Import sklearn components
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report
from sklearn.decomposition import NMF, TruncatedSVD

# Suppress warnings
import warnings
warnings.filterwarnings('ignore')
```

## 1.2 1. Data Preparation

### 1.2.1 1.1 Loading and Splitting Data

```
[22]: # Load datasets
train_df = pd.read_csv('../data/BBC News Train.csv', names=['ArticleId', 'Text', 'Category'], header=0)
kaggle_test_df = pd.read_csv('../data/BBC News Test.csv', names=['ArticleId', 'Text'], header=0)

# Split training data into train/validation sets for proper evaluation
train_texts, val_texts, train_labels, val_labels = train_test_split(
    train_df['Text'], train_df['Category'],
    test_size=0.2, random_state=42,
    stratify=train_df['Category']
)

print(f"Full training set shape: {train_df.shape}")
print(f"Training texts: {len(train_texts)}")
print(f"Validation texts: {len(val_texts)}")
print(f"Kaggle test set shape: {kaggle_test_df.shape}")

# Display class distribution
print("\nClass distribution in full training set:")
print(train_df['Category'].value_counts())
```

Full training set shape: (1490, 3)

Training texts: 1192

Validation texts: 298

Kaggle test set shape: (735, 2)

Class distribution in full training set:

Category

sport 346

business 336

politics 274

entertainment 273

tech 261

Name: count, dtype: int64

## 1.3 2. Model Implementation

### 1.3.1 2.1 Helper Functions

```
[23]: def prepare_data(train_texts, val_texts, kaggle_test_texts=None):
    # Initialize TF-IDF vectorizer
    tfidf = TfidfVectorizer(max_features=5000, stop_words='english')

    # Fit and transform data
```

```

X_train = tfidf.fit_transform(train_texts)
X_val = tfidf.transform(val_texts)

if kaggle_test_texts is not None:
    X_kaggle = tfidf.transform(kaggle_test_texts)
    return X_train, X_val, X_kaggle, tfidf

return X_train, X_val, tfidf

def train_evaluate_model(model, X_train, y_train, X_val, y_val, model_name):
    # Train and evaluate
    model.fit(X_train, y_train)
    train_pred = model.predict(X_train)
    val_pred = model.predict(X_val)

    # Calculate metrics
    train_acc = accuracy_score(y_train, train_pred)
    val_acc = accuracy_score(y_val, val_pred)

    # Calculate confusion matrix
    cm = confusion_matrix(y_val, val_pred)

    print(f"\n{model_name} Results:")
    print("-" * 20)
    print(f"Training Accuracy: {train_acc:.3f}")
    print(f"Validation Accuracy: {val_acc:.3f}")
    print("\nClassification Report:")
    print(classification_report(y_val, val_pred))

    return {
        'train_acc': train_acc,
        'val_acc': val_acc,
        'confusion_matrix': cm
    }

def evaluate_unsupervised_model(model, X_train, y_train, X_val, y_val):
    """Evaluate unsupervised model with proper topic-to-category mapping."""
    try:
        # Transform data
        train_topics = model.fit_transform(X_train)
        val_topics = model.transform(X_val)

        # Get dominant topic for each document
        train_doc_topics = train_topics.argmax(axis=1)
        val_doc_topics = val_topics.argmax(axis=1)

        # Count category occurrences for each topic

```

```

topic_counts = defaultdict(lambda: defaultdict(int))
for topic, category in zip(train_doc_topics, y_train):
    topic_counts[topic][category] += 1

# Get most common category overall as default
default_category = pd.Series(y_train).value_counts().index[0]

# Map topics to categories
topic_mapping = {}
for topic in range(train_topics.shape[1]):
    if topic in topic_counts and topic_counts[topic]:
        # Get category with highest count for this topic
        topic_mapping[topic] = max(topic_counts[topic].items(),
    ↪key=lambda x: x[1])[0]
    else:
        # Use default category if topic has no documents
        topic_mapping[topic] = default_category

# Make predictions
train_pred = [topic_mapping[topic] for topic in train_doc_topics]
val_pred = [topic_mapping[topic] for topic in val_doc_topics]

# Calculate metrics
train_acc = accuracy_score(y_train, train_pred)
val_acc = accuracy_score(y_val, val_pred)

return train_acc, val_acc

except Exception as e:
    print(f"Error in unsupervised evaluation: {str(e)}")
    return 0.0, 0.0

```

## 1.4 3. Data Efficiency Study

### 1.4.1 3.1 Training Size Experiments

```

[24]: # Prepare full dataset
X_train_full, X_val_full, X_kaggle, tfidf = prepare_data(train_texts,
    ↪val_texts, kaggle_test_df['Text'])
y_train_full = train_labels
y_val_full = val_labels

print("Conducting data efficiency study with different training set sizes...")

# Initialize supervised models
supervised_models = {
    'Naive Bayes': MultinomialNB(),

```

```

    'Logistic Regression': LogisticRegression(max_iter=1000, random_state=42)
}

# Test different training set sizes
train_sizes = [0.1, 0.2, 0.5, 1.0]
supervised_results = []

for size in train_sizes:
    print(f"\nTraining with {size*100}% of data")
    print("-" * 30)

    if size < 1.0:
        try:
            # Sample subset of training data
            train_subset, _, y_train_subset, _ = train_test_split(
                train_texts, train_labels,
                train_size=size, random_state=42,
                stratify=train_labels if size >= 0.1 else None
            )
        except ValueError:
            # If stratification fails, try without it
            train_subset, _, y_train_subset, _ = train_test_split(
                train_texts, train_labels,
                train_size=size, random_state=42
            )
        X_train_subset, X_val_subset, _ = prepare_data(train_subset, val_texts)
    else:
        X_train_subset = X_train_full
        y_train_subset = y_train_full
        X_val_subset = X_val_full

    for name, model in supervised_models.items():
        result = train_evaluate_model(
            model, X_train_subset, y_train_subset,
            X_val_subset, y_val_full, name
        )

        supervised_results.append({
            'Model': name,
            'Training Size': f"{size*100}%",
            'Train Accuracy': result['train_acc'],
            'Validation Accuracy': result['val_acc']
        })

# Create summary DataFrame
results_df = pd.DataFrame(supervised_results)

```

Conducting data efficiency study with different training set sizes...



Training with 10.0% of data

Naive Bayes Results:

Training Accuracy: 1.000

Validation Accuracy: 0.883

Classification Report:

	precision	recall	f1-score	support
business	0.76	0.96	0.85	67
entertainment	0.96	0.89	0.92	55
politics	0.94	0.85	0.90	55
sport	0.87	0.99	0.93	69
tech	1.00	0.67	0.80	52
accuracy			0.88	298
macro avg	0.91	0.87	0.88	298
weighted avg	0.90	0.88	0.88	298

Logistic Regression Results:

Training Accuracy: 1.000

Validation Accuracy: 0.886

Classification Report:

	precision	recall	f1-score	support
business	0.78	0.94	0.85	67
entertainment	0.93	0.93	0.93	55
politics	0.94	0.80	0.86	55
sport	0.89	0.99	0.94	69
tech	0.97	0.73	0.84	52
accuracy			0.89	298
macro avg	0.90	0.88	0.88	298
weighted avg	0.90	0.89	0.88	298

Training with 20.0% of data

Naive Bayes Results:

Training Accuracy: 1.000

Validation Accuracy: 0.926

Classification Report:

	precision	recall	f1-score	support
business	0.85	0.96	0.90	67
entertainment	0.96	0.96	0.96	55
politics	0.94	0.87	0.91	55
sport	0.92	1.00	0.96	69
tech	1.00	0.81	0.89	52
accuracy			0.93	298
macro avg	0.94	0.92	0.92	298
weighted avg	0.93	0.93	0.93	298

Logistic Regression Results:

-----

Training Accuracy: 1.000

Validation Accuracy: 0.930

Classification Report:

	precision	recall	f1-score	support
business	0.88	0.96	0.91	67
entertainment	0.95	0.96	0.95	55
politics	0.94	0.87	0.91	55
sport	0.93	1.00	0.97	69
tech	0.98	0.83	0.90	52
accuracy			0.93	298
macro avg	0.93	0.92	0.93	298
weighted avg	0.93	0.93	0.93	298

Training with 50.0% of data

-----

Naive Bayes Results:

-----

Training Accuracy: 0.997

Validation Accuracy: 0.953

Classification Report:

	precision	recall	f1-score	support
business	0.92	0.97	0.94	67
entertainment	0.98	0.98	0.98	55

politics	0.94	0.91	0.93	55
sport	0.97	1.00	0.99	69
tech	0.96	0.88	0.92	52
accuracy			0.95	298
macro avg	0.95	0.95	0.95	298
weighted avg	0.95	0.95	0.95	298

#### Logistic Regression Results:

-----

Training Accuracy: 1.000

Validation Accuracy: 0.956

#### Classification Report:

	precision	recall	f1-score	support
business	0.92	0.97	0.94	67
entertainment	0.96	0.98	0.97	55
politics	0.94	0.93	0.94	55
sport	1.00	1.00	1.00	69
tech	0.96	0.88	0.92	52
accuracy			0.96	298
macro avg	0.96	0.95	0.95	298
weighted avg	0.96	0.96	0.96	298

Training with 100.0% of data

-----

#### Naive Bayes Results:

-----

Training Accuracy: 0.992

Validation Accuracy: 0.977

#### Classification Report:

	precision	recall	f1-score	support
business	0.96	0.97	0.96	67
entertainment	1.00	1.00	1.00	55
politics	0.96	0.95	0.95	55
sport	1.00	1.00	1.00	69
tech	0.96	0.96	0.96	52
accuracy			0.98	298
macro avg	0.98	0.98	0.98	298
weighted avg	0.98	0.98	0.98	298

Logistic Regression Results:

-----

Training Accuracy: 0.997

Validation Accuracy: 0.966

Classification Report:

	precision	recall	f1-score	support
business	0.94	0.97	0.96	67
entertainment	0.96	1.00	0.98	55
politics	0.98	0.93	0.95	55
sport	0.99	1.00	0.99	69
tech	0.96	0.92	0.94	52
accuracy			0.97	298
macro avg	0.97	0.96	0.97	298
weighted avg	0.97	0.97	0.97	298

## 1.5 4. Comparison with Matrix Factorization

### 1.5.1 4.1 Unsupervised Model Evaluation

```
[25]: print("\nComparing with Matrix Factorization results...")

# Initialize unsupervised models with best parameters from notebook 2
unsupervised_models = {
    'NMF': NMF(n_components=5, random_state=42),
    'SVD': TruncatedSVD(n_components=5, random_state=42)
}

unsupervised_results = []

for size in train_sizes:
    if size < 1.0:
        try:
            # Sample subset of training data
            train_subset, _, y_train_subset, _ = train_test_split(
                train_texts, train_labels,
                train_size=size, random_state=42,
                stratify=train_labels if size >= 0.1 else None
            )
        except ValueError:
            # If stratification fails, try without it
            train_subset, _, y_train_subset, _ = train_test_split(
                train_texts, train_labels,
```

```

        train_size=size, random_state=42
    )
    X_train_subset, X_val_subset, _ = prepare_data(train_subset, val_texts)
else:
    X_train_subset = X_train_full
    y_train_subset = y_train_full
    X_val_subset = X_val_full

for name, model in unsupervised_models.items():
    # Evaluate unsupervised model with proper topic mapping
    train_acc, val_acc = evaluate_unsupervised_model(
        model, X_train_subset, y_train_subset,
        X_val_subset, y_val_full
    )

    print(f"\n{name} Results with {size*100}% data:")
    print("-" * 20)
    print(f"Training Accuracy: {train_acc:.3f}")
    print(f"Validation Accuracy: {val_acc:.3f}")

    unsupervised_results.append({
        'Model': name,
        'Training Size': f"{size*100}%",
        'Train Accuracy': train_acc,
        'Validation Accuracy': val_acc
    })

# Add unsupervised results to DataFrame
results_df = pd.concat([
    results_df,
    pd.DataFrame(unsupervised_results)
])

```

Comparing with Matrix Factorization results...

NMF Results with 10.0% data:

-----

Training Accuracy: 0.899

Validation Accuracy: 0.896

SVD Results with 10.0% data:

-----

Training Accuracy: 0.445

Validation Accuracy: 0.309

NMF Results with 20.0% data:

-----

Training Accuracy: 0.937  
Validation Accuracy: 0.916

SVD Results with 20.0% data:

-----

Training Accuracy: 0.487  
Validation Accuracy: 0.329

NMF Results with 50.0% data:

-----

Training Accuracy: 0.901  
Validation Accuracy: 0.903

SVD Results with 50.0% data:

-----

Training Accuracy: 0.461  
Validation Accuracy: 0.396

NMF Results with 100.0% data:

-----

Training Accuracy: 0.914  
Validation Accuracy: 0.913

SVD Results with 100.0% data:

-----

Training Accuracy: 0.439  
Validation Accuracy: 0.393

## 1.6 5. Results Visualization

```
[26]: # Plot learning curves
plt.figure(figsize=(15, 10))

# Plot 1: Test Accuracy vs Training Size
plt.subplot(2, 1, 1)
for model in results_df['Model'].unique():
    model_results = results_df[results_df['Model'] == model]
    plt.plot(
        model_results['Training Size'],
        model_results['Validation Accuracy'],
        marker='o',
        label=model
    )

plt.title('Validation Accuracy vs Training Data Size')
plt.xlabel('Training Data Size')
plt.ylabel('Validation Accuracy')
```

```

plt.legend()
plt.grid(True)


# Plot 2: Overfitting Analysis (Train vs Validation Accuracy)
plt.subplot(2, 1, 2)
for model in results_df['Model'].unique():
    model_results = results_df[results_df['Model'] == model]
    plt.plot(
        model_results['Training Size'],
        model_results['Train Accuracy'] - model_results['Validation Accuracy'],
        marker='o',
        label=model
    )

plt.title('Overfitting Analysis (Train - Validation Accuracy)')
plt.xlabel('Training Data Size')
plt.ylabel('Accuracy Gap (Train - Validation)')
plt.legend()
plt.grid(True)

plt.tight_layout()
plt.show()

# Display summary table
print("\nPerformance Summary:")
summary_df = results_df.pivot_table(
    index=['Model', 'Training Size'],
    values=['Train Accuracy', 'Validation Accuracy'],
    aggfunc='first'
).round(3)
print(summary_df)

```



### Performance Summary:

		Train Accuracy		Validation Accuracy
Model	Training Size			
Logistic Regression	10.0%	1.000		0.886
	100.0%	0.997		0.966
	20.0%	1.000		0.930
	50.0%	1.000		0.956
NMF	10.0%	0.899		0.896
	100.0%	0.914		0.913
	20.0%	0.937		0.916
	50.0%	0.901		0.903
Naive Bayes	10.0%	1.000		0.883
	100.0%	0.992		0.977
	20.0%	1.000		0.926
	50.0%	0.997		0.953
SVD	10.0%	0.445		0.309
	100.0%	0.439		0.393
	20.0%	0.487		0.329
	50.0%	0.461		0.396



## 1.7 6. Analysis and Conclusions

### 1.7.1 6.1 Data Efficiency Analysis

Training size impact by model: 1. Naive Bayes shows strong performance with limited data 2. Logistic Regression requires more data for optimal results 3. Unsupervised methods need larger datasets for stability

### 1.7.2 6.2 Overfitting Analysis

Model stability characteristics: 1. Naive Bayes shows minimal overfitting 2. Logistic Regression exhibits moderate overfitting 3. Matrix factorization methods show less overfitting but lower accuracy

### 1.7.3 6.3 Trade-offs Analysis

**Supervised Approaches Advantages:** 1. Higher accuracy across all training sizes 2. Better performance with limited data 3. More consistent results

**Disadvantages:** 1. Require labeled data 2. Show more overfitting 3. May not generalize to new categories

**Unsupervised Approaches Advantages:** 1. No need for labeled data 2. Less overfitting 3. Can discover latent patterns

**Disadvantages:** 1. Lower overall accuracy 2. Require more data for stable results 3. Less interpretable results

### 1.7.4 6.4 Recommendations

1. For optimal accuracy:
  - Use Logistic Regression with full dataset
  - Consider ensemble methods
  - Focus on feature engineering
2. For limited data scenarios:
  - Prefer supervised methods
  - Use Naive Bayes for better generalization
  - Focus on feature selection
3. For unlabeled data:
  - Start with matrix factorization
  - Use larger training sets
  - Consider semi-supervised approaches